

Nachos Project 3

B08505048 工海四 劉名凱

1. Motivation

Unlike project 1 and 2, task in project 3 require much more memory space. Therefore, we need to implement page replacement, swap out/in frames to acquire enough memory space and thus enable program finish successfully. Besides, project 3 also require multi-programming implemented in project 1. Only combining multi-programming and page replacement can we tackle down project 3.

The first thing to do is to decide where to deal with page fault. I choose to handle page fault whenever a page is needed and is not loaded into memory, avoiding producing any `pageFaultException`. Page replacement also need a disk to swap in/out frames. I follow the hint in the TA's pdf, using `Synchdisk` as a virtual disk. Finally, we only need to implement how page was replaced when a page fault occurs.

2. Implementation

First, add a `SynchDisk` in `userkernel.h` to act as a virtual disk, so that we can swap frames into/from the disk.

```
20 class UserProgKernel : public ThreadedKernel {
21 public:
22     UserProgKernel(int argc, char **argv);
23     // Interpret command line arguments
24     ~UserProgKernel(); // deallocate the kernel
25
26     void Initialize(); // initialize the kernel
27
28     void Run(); // do kernel stuff
29
30     void SelfTest(); // test whether kernel is working
31
32     SynchDisk *virtualMemoryDisk; // Daniel added
```

Also remember to initialize the `SynchDisk` in `userkernel.cc`

```
54 void
55 UserProgKernel::Initialize()
56 {
57     ThreadedKernel::Initialize(); // init multithreading
58
59     machine = new Machine(debugUserProg);
60     fileSystem = new FileSystem();
61     virtualMemoryDisk = new SynchDisk("Virtual Memory");
62 #ifdef FILESYS
63     synchDisk = new SynchDisk("New SynchDisk");
64 #endif // FILESYS
65 }
```

Add variable in machine.h:

- usedPhyPage: Record the usage of main memory
- usedVirPage: Record the usage of SynchDisk (virtual disk)
- FreePages: Record how many free frames in main memory
- frameTable: Record each frame's (in main memory) corresponding page (in each thread)
- ID_num: Record how many threads there are in one process
- fifo: First in first out queue, recording the sequence during allocating frames in main memory

```
89 class Machine {
90 public:
91     Machine(bool debug); // Initialize the simulation of the hardware
92     // for running user programs
93     ~Machine(); // De-allocate the data structures
94
95 // Routines callable by the Nachos kernel
96 void Run(); // Run a user program
97
98 int ReadRegister(int num); // read the contents of a CPU register
99
100 void WriteRegister(int num, int value);
101 // store a value into a CPU register
102
103 // Daniel added
104 bool usedPhyPage[NumPhysPages];
105 bool usedVirPage[NumPhysPages];
106 int FreePages;
107 TranslationEntry *frameTable[NumPhysPages];
108 int ID_num;
109 std::queue<int> fifo;
```

Initialize variables in machine.cc

```
54 Machine::Machine(bool debug)
55 {
56     // Daniel added
57     ID_num = 0;
58     FreePages = NumPhysPages;
59     for(unsigned int i = 0; i<NumPhysPages; i++){
60         usedPhyPage[i] = false;
61         usedVirPage[i] = false;
62     }
63 }
```

Add variable in addrspace.h:

ID: Thread's id

pageTableloaded: check whether pageTable is fully loaded

```
22 class AddrSpace {
23 public:
24     //
25     int ID;
26     AddrSpace();      // Create an address space.
27     ~AddrSpace();     // De-allocate an address space
28
29     void Execute(char *fileName); // Run the the program
30     // stored in the file "executable"
31
32     void SaveState();    // Save/restore address space-specific
33     void RestoreState(); // info on a context switch
34
35 private:
36     bool pageTableloaded;
37     TranslationEntry *pageTable; // Assume linear page table translation
38     // for now!
39     unsigned int numPages;    // Number of pages in the virtual
40     // address space
41
42     bool Load(char *fileName); // Load the program into memory
43     // return false if not found
44
45     void InitRegisters(); // Initialize user-level CPU registers,
46     // before jumping to user code
47 };
48
```

Next step is to make sure when loading thread's address space, thread can run successfully even when main memory is full. To do this, we need to modify code in `addrspace::Load` which control the process of loading address space.

1. Remove `ASSERTION(numPages <= FreePages)` since we can deal with page fault now.

```
88 bool
89 AddrSpace::Load(char *fileName)
90 {
91     OpenFile *executable = kernel->fileSystem->Open(fileName);
92     NoffHeader noffH;
93     unsigned int size;
94
95     if (executable == NULL) {
96         cerr << "Unable to open file " << fileName << "\n";
97         return FALSE;
98     }
99     executable->ReadAt((char *)&noffH, sizeof(noffH), 0);
100     if ((noffH.noffMagic != NOFFMAGIC) && (WordToHost(noffH.noffMagic) == NOFFMAGIC))
101         SwapHeader(&noffH);
102     ASSERT(noffH.noffMagic == NOFFMAGIC);
103
104     // how big is address space?
105     size = noffH.code.size + noffH.initData.size + noffH.uninitData.size
106         + UserStackSize; // we need to increase the size
107         // to leave room for the stack
108     numPages = divRoundUp(size, PageSize);
109     cout << "number of pages of " << fileName << " is " << numPages << endl;
110     size = numPages * PageSize;
111
```

2. Initializing pageTable and printing how many frames this thread requires.

```
111
112 // Daniel added
113 // load page table
114 pageTable = new TranslationEntry[numPages];
115 for (unsigned int i = 0; i < numPages; i++) {
116     pageTable[i].virtualPage = i;
117     pageTable[i].physicalPage = i;
118     pageTable[i].valid = true;
119     pageTable[i].use = false;
120     pageTable[i].dirty = false;
121     pageTable[i].readOnly = false;
122 }
123 printf("Thread %d requires %d frame, %d free frame remain\n", AddrSpace::ID, numPages, kernel->machine->FreePages);
124 DEBUG(dbgAddr, "Initializing address space: " << numPages << ", " << size);
```

3. Allocating frames to thread. Looping through all pages this thread needs. If main memory has space, assigning frame to page. Make sure to record frame's corresponding page and push frame number into fifo. Then read file's content according to thread's allocated main memory size.

```
125
126 // then, copy in the code and data segments into memory
127 if (noffH.code.size > 0) {
128     DEBUG(dbgAddr, "Initializing code segment.");
129     DEBUG(dbgAddr, noffH.code.virtualAddr << ", " << noffH.code.size);
130     // Daniel added
131     // allocate main memory
132     for(unsigned int j = 0, i = 0; i < numPages; i++){
133         j = 0;
134         while(j < NumPhysPages && kernel->machine->usedPhyPage[j] == true)j++;
135         if(j < NumPhysPages) {
136             bzero(&kernel->machine->mainMemory[j * PageSize], PageSize);
137             kernel->machine->usedPhyPage[j] = true;
138             kernel->machine->FreePages--;
139             kernel->machine->frameTable[j] = &pageTable[i];
140             pageTable[i].physicalPage = j;
141             pageTable[i].valid = true;
142             pageTable[i].use = false;
143             pageTable[i].dirty = false;
144             pageTable[i].readOnly = false;
145             pageTable[i].ID = ID;
146             // Add physical address to FIFO queue
147             kernel->machine->fifo.push(j);
148
149             executable->ReadAt( &(kernel->machine->mainMemory[j * PageSize]), PageSize,
150                 noffH.code.inFileAddr + (i*PageSize));
151         }
152     }
```

4. If main memory's space is not enough, find space in SynchDisk (using first fit) and read file's content into SynchDisk. Make sure to set pageTable[i].valid = false, so that kernel knows this page is not in main memory.

```
152 // requires virtual memory
153 else {
154     char *buffer;
155     buffer = new char[PageSize];
156     j = 0;
157     // find next empty virtual page
158     while(kernel->machine->usedVirPage[j] != false){ j++; }
159
160     kernel->machine->usedVirPage[j]=true;
161     pageTable[i].virtualPage = j;
162     pageTable[i].valid = false;
163     pageTable[i].use = false;
164     pageTable[i].dirty = false;
165     pageTable[i].readOnly = false;
166     executable->ReadAt(buffer, PageSize, noffH.code.inFileAddr + (i * PageSize));
167     kernel->virtualMemoryDisk->WriteSector(j, buffer);
168 }
```


5. Set `pageTableloaded = true` after `pageTable` is successful loaded. (`RestoreState` will load the pagetable).

```
190 void
191 AddrSpace::Execute(char *fileName)
192 {
193     pageTableloaded = false;
194     if (!Load(fileName)) {
195         cout << "inside !Load(fileName)" << endl;
196         return; // executable not found
197     }
198
199     //kernel->currentThread->space = this;
200     this->InitRegisters(); // set the initial register values
201     this->RestoreState(); // load page table register
202     pageTableloaded = true;
203     kernel->machine->Run(); // jump to the user program
204     ASSERTNOTREACHED(); // machine->Run never returns;
205     // the address space exits
206     // by doing the syscall "exit"
207 }
```

6. On a context switch, save `pageTable` status only if `pageTable` is successful loaded.

```
251 void AddrSpace::SaveState()
252 {
253     if (pageTableloaded){
254         pageTable=kernel->machine->pageTable;
255         numPages=kernel->machine->pageTableSize;
256     }
257 }
```

After finish modifying the process of loading thread's address space. We need to implement page replacement. I implement page replacement in `translate.cc`. There are two cases when a page fault occurred

Case1: Some frames are released, so we can just allocate that frame to thread.

```
213 } else if (!pageTable[vpn].valid) { Page is not in main memory, page fault occurs
214     DEBUG(dbgAddr, "Invalid virtual page # " << virtAddr);
215     // Daniel added
216     kernel->stats->numPageFaults++;
217     unsigned int j = 0; Find is there any frames that are released
218     while(kernel->machine->usedPhyPage[j] != false && j< NumPhysPages) j++;
219     if(j < NumPhysPages){ If there are available frame
220         char *buffer;
221         buffer = new char[PageSize];
222         kernel->machine->usedPhyPage[j] = true;
223         kernel->machine->FreePages--;
224         kernel->machine->frameTable[j] = &pageTable[vpn];
225         pageTable[vpn].physicalPage = j;
226         pageTable[vpn].valid = true;
227
228         kernel->virtualMemoryDisk->ReadSector(pageTable[vpn].virtualPage, buffer);
229         bcopy(buffer, &mainMemory[j * PageSize], PageSize);
230     }
```

Case 2: There are no available frame, swap out a frame (using first in first out) and swap in page's corresponding content from SynchDisk.

```
231 else{
232     char *buffer1;
233     buffer1 = new char[PageSize];
234     char *buffer2;
235     buffer2 = new char[PageSize];
236
237     //FIFO
238     int victim = fifo.front();
239     fifo.pop();
240
241     // 把東西存到 disk 裡面
242     bcopy(&mainMemory[victim * PageSize], buffer1, PageSize); // copy victim content to buffer 1
243     kernel->virtualMemoryDisk->ReadSector(pageTable[vpn].virtualPage, buffer2); // swap vpn content into buffer 2
244     bcopy(buffer2, &mainMemory[victim*PageSize], PageSize); // copy vpn content into main memory
245     kernel->virtualMemoryDisk->WriteSector(pageTable[vpn].virtualPage, buffer1); // swap victim content into disk
246
247
248     kernel->machine->frameTable[victim]->virtualPage = pageTable[vpn].virtualPage;
249     kernel->machine->frameTable[victim]->valid = false;
250
251     // 把東西 load 到 main memory 中
252     pageTable[vpn].valid = true;
253     pageTable[vpn].physicalPage = victim;
254     kernel->machine->frameTable[victim] = &pageTable[vpn];
255
256     fifo.push(victim);
257 }
```

3. Result

matmult.c

```
daniel@daniel-VirtualBox: ~/downloads/nachos-4.0/code/userprog
daniel@daniel-VirtualBox:~/downloads/nachos-4.0/code/userprog$ ./nachos -e ../test/matmult
RR
Total threads number is 1
Thread ../test/matmult is executing.
Thread 0 requires 54 frame, 64 free frame remain
return value:7220
No threads ready or runnable, and no pending interrupts.
Assuming the program completed.
Machine halting!

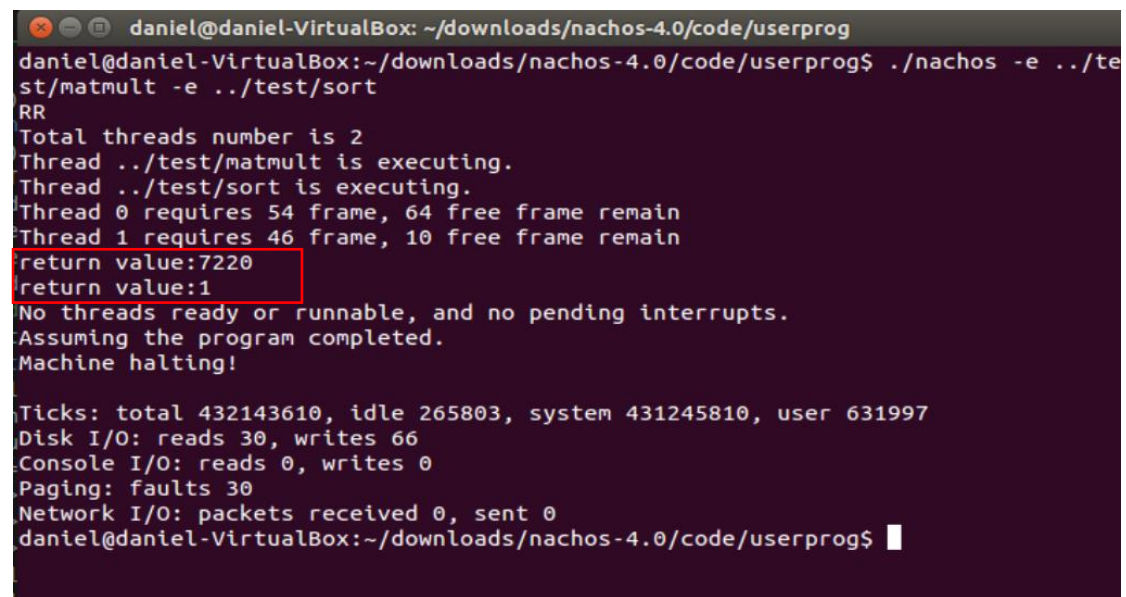
Ticks: total 702300, idle 59, system 70250, user 631991
Disk I/O: reads 0, writes 0
Console I/O: reads 0, writes 0
Paging: faults 0
Network I/O: packets received 0, sent 0
daniel@daniel-VirtualBox:~/downloads/nachos-4.0/code/userprog$
```

sort.c

```
daniel@daniel-VirtualBox: ~/downloads/nachos-4.0/code/userprog
daniel@daniel-VirtualBox:~/downloads/nachos-4.0/code/userprog$ ./nachos -e ../test/sort
RR
Total threads number is 1
Thread ../test/sort is executing.
Thread 0 requires 46 frame, 64 free frame remain
return value:1
No threads ready or runnable, and no pending interrupts.
Assuming the program completed.
Machine halting!

Ticks: total 43117300, idle 38, system 4311750, user 38805512
Disk I/O: reads 0, writes 0
Console I/O: reads 0, writes 0
Paging: faults 0
Network I/O: packets received 0, sent 0
daniel@daniel-VirtualBox:~/downloads/nachos-4.0/code/userprog$
```

matmult.c and sort.c simultaneously

A terminal window titled 'daniel@daniel-VirtualBox: ~/downloads/nachos-4.0/code/userprog'. The prompt is 'daniel@daniel-VirtualBox:~/downloads/nachos-4.0/code/userprog\$'. The user has entered './nachos -e ../test/matmult -e ../test/sort'. The output shows 'RR', 'Total threads number is 2', 'Thread ../test/matmult is executing.', 'Thread ../test/sort is executing.', 'Thread 0 requires 54 frame, 64 free frame remain', 'Thread 1 requires 46 frame, 10 free frame remain', 'return value:7220', 'return value:1', 'No threads ready or runnable, and no pending interrupts.', 'Assuming the program completed.', 'Machine halting!', and various statistics: 'Ticks: total 432143610, idle 265803, system 431245810, user 631997', 'Disk I/O: reads 30, writes 66', 'Console I/O: reads 0, writes 0', 'Paging: faults 30', 'Network I/O: packets received 0, sent 0'. The prompt is now 'daniel@daniel-VirtualBox:~/downloads/nachos-4.0/code/userprog\$' with a cursor. A red box highlights the two 'return value' lines.

```
daniel@daniel-VirtualBox: ~/downloads/nachos-4.0/code/userprog
daniel@daniel-VirtualBox:~/downloads/nachos-4.0/code/userprog$ ./nachos -e ../te
st/matmult -e ../test/sort
RR
Total threads number is 2
Thread ../test/matmult is executing.
Thread ../test/sort is executing.
Thread 0 requires 54 frame, 64 free frame remain
Thread 1 requires 46 frame, 10 free frame remain
return value:7220
return value:1
No threads ready or runnable, and no pending interrupts.
Assuming the program completed.
Machine halting!

Ticks: total 432143610, idle 265803, system 431245810, user 631997
Disk I/O: reads 30, writes 66
Console I/O: reads 0, writes 0
Paging: faults 30
Network I/O: packets received 0, sent 0
daniel@daniel-VirtualBox:~/downloads/nachos-4.0/code/userprog$
```

At first, I couldn't successfully run mamult.c and sort.c at the same time. Kernel kept reporting illegal instruction exception. After debugging for three days, I found out that I forgot to load file into SynchDisk (virtual disk). Therefore, nothing was in the disk when swapping occurred.

There is another page replacement method called Least Recently Used which swap out least recently used page. All I need to do is to add a variable in pageTable to record number of each page is used. However, I didn't implement LRU, since FIFO is quick enough.